

This document is a complete, detailed technical thesis prepared for academic defense. It covers every phase of ZARI.ai from scratch: problem context, exploratory data analysis, each model trained (with full epoch-by-epoch error curves), every engineering hurdle we encountered and solved, per-class evaluation results, and the full multilingual Retrieval-Augmented Generation (RAG) advisory system.

Written so that any reader - even without prior knowledge - can fully understand the motivation, methodology, decisions, and final outcomes of ZARI.ai.

Quick Stats

Dataset:	49,805 images 26 classes 3 crops
Models:	Model A: EfficientNetV2-B2 Crop Router (99.48% Acc) Model B: 3x EDL Classifiers
Best F1:	Tomato 0.9787 Potato 0.9718 Pepper 0.9963
RAG:	208 multilingual chunks 384d embeddings Pashto score 0.5184
Latency:	12.05 ms end-to-end (CUDA) SCRC FAR = 1.04% Status: PRODUCTION_READY

Section 1: Problem Context & Motivation

1.1 The Agricultural Challenge in Pakistan

Pakistan's agricultural sector faces a chronic challenge: smallholder farmers who grow Tomato, Potato, and Bell Pepper lose 20-40% of annual yield to plant diseases. Losses are not just from disease itself, but from incorrect treatment - farmers spray broad-spectrum fungicides blindly, targeting the wrong pathogen, damaging the crop, and incurring unnecessary costs.

THREE ROOT PROBLEMS ZARI.ai WAS BUILT TO SOLVE:

PROBLEM 1 - MISDIAGNOSIS: Field workers cannot reliably distinguish visually similar diseases. For example, Tomato Early Blight (*Alternaria solani*) and Tomato Target Spot (*Corynespora cassiicola*) produce nearly identical concentric lesion patterns on leaves. Treating the wrong disease wastes money and worsens the outbreak.

PROBLEM 2 - OVERCONFIDENT AI SYSTEMS: Standard deep learning classifiers using Softmax output layers return high-confidence predictions even when shown images they have never been trained on. Showing a non-plant image still returns "98% Tomato Bacterial Spot" - a dangerously wrong and overconfident answer.

PROBLEM 3 - LANGUAGE BARRIERS: Plant disease resources are published in English. Rural farming communities in Khyber Pakhtunkhwa primarily speak Pashto, and the broader rural population uses Urdu. No single system existed that could answer agricultural queries in all three languages simultaneously.

1.2 ZARI.ai System Overview

- * Vision Diagnostic Engine: A two-stage EfficientNetV2-B2 hierarchical classifier. Model A identifies the crop type, then the matching Model B diagnoses the specific disease.
- * Safety Gate (SCRC): Selective Classification & Risk Control automatically rejects ambiguous or out-of-distribution images. False Acceptance Rate = 1.04%.
- * Visual Severity Proxy: Meta SAM2 leaf segmentation + Grad-CAM heatmap intersection. Estimates disease coverage (Mild/Moderate/Severe) in 4.57 ms.
- * Multilingual RAG Advisory: ChromaDB vector store (208 evidence chunks) queried in English, Urdu, and Pashto.

* IPM Compliance: Strict Integrated Pest Management hierarchy - Cultural -> Biological -> Chemical active ingredients only.

Section 2: Dataset Construction & Exploratory Data Analysis (EDA)

2.1 Raw Data Sources & The Challenge of Curation

We did not start with a clean dataset. Raw archives came from six repositories with different naming conventions, resolutions, and quality levels - approximately 68,000 images across 67 label categories with duplicates and incorrect labels.

Raw Sources: PlantVillage (Kaggle), Mendeley Plant Disease datasets, Pakistan local field samples, Bangladesh open dataset, PLD (Plant Leaf Disease) open dataset.

2.2 Data Cleaning Pipeline - 5 Steps

Step 1 - Exact Duplicate Removal

MD5 checksum hashing of every image. Identical MD5 hashes = exact binary duplicates, removed completely.

Step 2 - Near-Duplicate Removal (pHash)

Perceptual Hash (pHash) computed per image. Hamming distance ≤ 2 out of 64 bits = visually near-identical frames, deduplicated to one representative.

Step 3 - Taxonomy Consolidation

Mapped 67 raw label strings to 26 canonical disease class names. Required manually writing class_aliases_v3.yaml mapping file.

Step 4 - Quality Filtering

Removed: Laplacian blur score < 50 (too blurry), brightness < 30 or > 220 (under/overexposed), resolution $< 64 \times 64$ pixels.

Step 5 - Stratified 80/10/10 Split

Split into Train/Val/Test using stratified sampling - each class represented proportionally in all splits.

2.3 Final Master V4 Dataset Statistics

Total Images:	49,805 (after cleaning from ~68,000 raw)
Train Set:	39,834 images (80%)
Validation Set:	4,978 images (10%)
Test Set:	4,993 images (10%)
Crops:	3: Tomato, Potato, Bell Pepper
Disease Classes:	26 canonical classes total

Crop	Disease Classes	# Classes
Tomato	Bacterial Spot, Early Blight, Late Blight, Leaf Mold, Septoria, Spider Mites, Target Spot, TYLCV, Mosaic Virus, Verticillium Wilt, Fusarium	13
Potato	Early Blight, Late Blight, Healthy	3
Bell Pepper	Bacterial Spot, Cercospora Leaf Spot, Leaf Curl, Nutrition Deficiency, Powdery Mildew, Healthy	6

2.4 EDA Visualizations




Figure: Fig 2.1 - Class-level image count across 26 canonical disease classes. Shows class imbalance: e.g., Tomato Healthy has significantly more samples than Tomato Mosaic Virus.




Figure: Fig 2.2 - Total sample volume per crop. Tomato ~60%, Pepper ~25%, Potato ~15%. Imbalance handled by per-crop stratified splits.




Figure: Fig 2.3 - Healthy vs Diseased ratio. Diseased samples dominate (~78%), reflecting real field disease prevalence.

5. Split Distribution




Figure: Fig 2.4 - Train / Validation / Test split counts per class. Confirms proportional stratified 80/10/10 split across all 26 classes.

8. Blur Score Distribution (Clipped)




Figure: Fig 2.5 - Laplacian blur score distribution. Long tail of blurry images (score < 50) removed during quality control step.




Figure: Fig 2.6 - Image brightness distribution. Images outside 30-220 range (underexposed / overexposed) excluded from dataset.




Figure: Fig 2.7 - Crop x Pathogen co-occurrence heatmap. Shows which pathogen types (Fungal, Bacterial, Viral, Nutritional) affect which crops.




Figure: Fig 2.8 - Estimated disease severity distribution. Mild, Moderate, Severe categories derived from visual leaf coverage analysis.

Section 3: System Architecture - 9-Stage End-to-End Pipeline

ZARI.ai's inference pipeline has 9 sequential stages. Each stage performs a specific function and passes structured data to the next. The entire pipeline completes in 12.05 ms on CUDA GPU.

Stage 1 - Image Acquisition & Preprocessing

Input image (JPEG/PNG) loaded, decoded, resized to 256x256 RGB pixels, converted to PyTorch float32 tensor, normalised using ImageNet mean=[0.485,0.456,0.406] and std=[0.229,0.224,0.225]. This normalisation is essential because EfficientNetV2-B2 was pretrained on ImageNet with these exact statistics - mismatched normalisation degrades accuracy.

Stage 2 - Model A: Crop Router Classification

EfficientNetV2-B2 (7.71M parameters) classifies the image into one of 3 crop categories: Tomato (class 0), Potato (class 1), Bell Pepper (class 2). This is a prerequisite - without knowing the crop, the correct Model B cannot be selected. Model A achieved 99.48% test accuracy and Macro F1 = 0.9926.

Stage 3 - Model B: EDL Disease Classifier

Based on Model A's crop output, the corresponding crop-specific Model B is selected (Tomato: 13 classes, Potato: 3 classes, Pepper: 6 classes). Each Model B uses an EfficientNetV2-B2 backbone with a Evidential Deep Learning (EDL) output head. Instead of Softmax, EDL outputs non-negative evidence scores $e_k = \text{Softplus}(z_k)$ for each class k , parameterising a Dirichlet distribution ($\alpha_k = e_k + 1$).

Stage 4 - SCRC Safety Gate

SCRC checks: (a) Model A confidence < 0.85, OR (b) Evidential Uncertainty $u = K/S > 0.45$. If EITHER triggers, pipeline returns REJECT with 'insufficient confidence for disease-specific recommendation' in English and Urdu. Prevents wrong diagnosis from reaching farmer. False Acceptance Rate = 1.04%.

Stage 5 - SAM2 Leaf Mask Segmentation

Meta SAM2 (Hiera-Tiny, 38.9M parameters) receives a central bounding-box prompt [10%W, 10%H, 90%W, 90%H]. SAM2 outputs a binary mask isolating the leaf from background soil, weeds, and pots. Completes in 4.57 ms on CUDA GPU.

Stage 6 - Grad-CAM Heatmap & Severity Calculation

Grad-CAM maps activation gradients at EfficientNetV2-B2 layer features.7.1. We intersect the SAM2 leaf mask with the Grad-CAM heatmap (activation >= 0.50). Fraction of leaf pixels with high activation = Visual Disease Coverage %. Coverage < 15% = Mild, 15-35% = Moderate, > 35% = Severe.

Stage 7 - Weather Context Integration

Ambient weather data (temperature, humidity, rainfall) injected into pipeline context. If severe visual coverage (> 35%) coincides with epidemic-risk humidity (> 85% RH), a COMBINED URGENCY WARNING flag is set to True.

Stage 8 - ChromaDB Multilingual Vector RAG Search

Query encoded into 384d dense vector by paraphrase-multilingual-MiniLM-L12-v2 (English, Urdu, or Pashto). Searched against ChromaDB HNSW-indexed collection 'zari_3crop_treatment_kb' (208 evidence chunks). Top-k most semantically similar chunks retrieved as advisory evidence.

Stage 9 - IPM Advisory Synthesis

Evidence chunks assembled into advisory following strict IPM hierarchy: Cultural Control -> Biological Control -> Chemical Control (active ingredients only: Mancozeb, Metalaxyl-M, Copper Hydroxide). No hallucinated trade names or unverified PHI wait days.

Section 4: Model A - Crop Router Training & Results

4.1 Architecture & Training Setup

Model A's sole job: look at an image and decide if it is Tomato, Potato, or Bell Pepper - a 3-class classification problem.

Architecture: EfficientNetV2-B2 (tf_efficientnetv2_b2) from the timm library, pretrained on ImageNet-21k then fine-tuned on our 3-crop dataset. Selected over heavier ViTs for its excellent accuracy/speed tradeoff with far fewer parameters.

Training Hyperparameters:

- Learning Rate: 3e-4 (Adam optimiser + cosine annealing LR schedule)
- Batch Size: 64 images per step
- Total Epochs: 10
- Loss Function: Cross-Entropy Loss
- Data Augmentation: Random horizontal/vertical flips, rotation +/-30 deg, colour jitter (brightness +/-0.3, contrast +/-0.3), random crop resize.
- Hardware: NVIDIA CUDA GPU (single GPU)

4.2 Epoch-by-Epoch Training Metrics

Epoch	Train Loss	Val Loss	Val Accuracy	Learning Rate
1	0.4464	0.2101	93.23%	0.000300
2	0.1393	0.1507	95.15%	0.000293
3	0.0888	0.1317	96.07%	0.000271
4	0.0596	0.1292	96.22%	0.000238
5	0.0384	0.1260	96.39%	0.000196
6	0.0250	0.1161	97.14%	0.000150
7	0.0153	0.1293	97.11%	0.000104
8	0.0105	0.1243	97.25%	0.000062
9	0.0081	0.1239	97.23%	0.000029
10	0.0069	0.1242	97.49%	0.000007

Interpretation: Training loss drops from 0.4464 (Epoch 1) to 0.0069 (Epoch 10) - the model learns the training set extremely effectively. Validation loss plateaus around 0.1242 by Epoch 10. Validation accuracy climbs from 93.23% to 97.49%. The learning rate decays from 3e-4 to 7e-6 following cosine annealing.

4.3 Training Curve Visualization




Figure: Fig 4.1 - Model A (Crop Router) training curves: Training Loss (blue), Validation Loss (orange), Validation Accuracy (green) over 10 epochs.

4.4 Final Test Set Results

Test Accuracy:	99.48%
Macro F1 Score:	0.9926
Parameters:	7,705,221 (7.71M)
File Size:	88.89 MB
Checkpoint:	ml_pipeline/checkpoints/model_a/best_model_a_efficientnetv2_b2.pth

Model A effectively routes 99.48% of images to the correct crop-specific classifier. The 0.52% error rate is handled by the SCRC safety gate in Stage 4.

Section 5: Evidential Deep Learning (EDL) - Theory & Model B Training

5.1 Why Standard Softmax Is Dangerous for Medical/Agricultural AI

Every standard classifier ends with a Softmax layer converting logits to probabilities that MUST sum to 1.0. This forces the network to always assign confidence to some class - even on inputs it has never seen.

CONCRETE EXAMPLE: Show a standard Softmax plant disease classifier a photo of a car. It will still output: "Tomato Bacterial Spot: 67%, Early Blight: 21%..." These are meaningless confidence values on a completely out-of-distribution input. In agriculture, this is dangerous: a farmer uploading a blurry or irrelevant photo could receive a confident but completely wrong disease diagnosis and spray unnecessary pesticides.

Evidential Deep Learning (EDL) SOLVES THIS by replacing Softmax with a Dirichlet distribution parameterised by learnable evidence values. When evidence is very low (unfamiliar input), the uncertainty score u automatically approaches 1.0, triggering a safe rejection.

5.2 The EDL Mathematical Framework

Given a K-class classification problem with input x :

STEP 1 - Raw logits z_k produced by the final fully connected layer.

STEP 2 - Evidence values: $e_k = \text{Softplus}(z_k) = \log(1 + \exp(z_k))$
Softplus guarantees non-negative evidence. Negative evidence is meaningless.

STEP 3 - Dirichlet parameters: $\alpha_k = e_k + 1$

The "+1" prior means even with zero evidence, uncertainty is bounded.

STEP 4 - Total Dirichlet Strength: $S = \sum(\alpha_k \text{ for } k = 1 \text{ to } K)$

High S = more total evidence = lower uncertainty.

STEP 5 - Expected class probability: $p_k = \alpha_k / S$

Class with highest p_k = predicted disease.

STEP 6 - Evidential Uncertainty: $u = K / S$

- Familiar image: model has high evidence -> S very large -> u approaches 0.
- Unfamiliar/OOD image: evidence stays near zero -> S near K -> u approaches 1.0.
- $u > 0.45$ triggers SCRC rejection.

TRAINING LOSS: EDL uses Type-II Maximum Likelihood loss (negative log-likelihood of Dirichlet over observed one-hot label) + KL-divergence regularisation penalising evidence for incorrect classes. This makes early epochs harder than standard CE.

5.3 Model B Training Setup

Three separate Model B classifiers trained, one per crop:

- Model B Tomato: 13-class EDL classifier
- Model B Potato: 3-class EDL classifier
- Model B Pepper: 6-class EDL classifier

Each uses EfficientNetV2-B2 backbone with modified classification head outputting raw logits fed through Softplus for EDL evidence values.

Training Hyperparameters (same for all 3):

- Learning Rate: $3e-4$ (Adam + cosine annealing)
- Batch Size: 64 per crop
- Total Epochs: 10
- Loss: EDL Type-II Max Likelihood + KL Divergence Regularisation

5.4 Epoch-by-Epoch Training Metrics (Combined Model B History)

Epoch	Train Loss	Val Loss	Val Accuracy	Mean Uncertainty
1	1.6092	1.3449	96.60%	0.6909
2	0.9147	1.0841	96.78%	0.5787
3	0.7478	0.9942	97.19%	0.5112
4	0.6502	0.9448	97.19%	0.4693
5	0.5820	0.9699	97.16%	0.4426
6	0.5329	1.0238	97.16%	0.4191
7	0.4919	1.0153	97.47%	0.4021
8	0.4633	1.0477	97.46%	0.3962
9	0.4472	1.0889	97.70%	0.3915
10	0.4415	1.1481	97.58%	0.3929

Interpretation: EDL training loss starts very high (1.6092) because the Dirichlet-based loss includes KL-divergence regularisation penalising unearned evidence. This makes first epochs harder. By Epoch 10, train loss reaches 0.4415. CRUCIALLY, Mean Uncertainty drops from 0.6909 (Epoch 1) to 0.3915 (Epoch 9) - the model becomes increasingly confident on familiar training samples as it learns the evidence distribution.

5.5 Training Curve Visualization




Figure: Fig 5.1 - Model B (EDL Disease Classifiers) training curves: loss convergence and uncertainty reduction over 10 epochs of Dirichlet evidence learning.

5.6 Final Test Set Results - Per Crop

Crop (# classes)	Test Accuracy	Macro F1	Test AUROC	Checkpoint File
Tomato (13 classes)	98.29%	0.9787	0.9993	best_model_b_tomato.pth
Potato (3 classes)	96.75%	0.9718	0.9963	best_model_b_potato.pth
Pepper (6 classes)	99.40%	0.9963	1.0000	best_model_b_pepper.pth

Section 6: Swin-Tiny Vision Transformer - Why We Evaluated and Rejected It

6.1 Why We Evaluated Swin-Tiny

Vision Transformers (ViTs) are state-of-the-art models using Self-Attention instead of convolutions. Swin-Tiny is a hierarchical ViT that processes images in non-overlapping local windows, making it computationally efficient.

We evaluated Swin-Tiny as a potential upgrade because:

- 1. ViTs show superior performance on several image classification benchmarks in literature.
- 2. Swin-Tiny demonstrated excellent accuracy on plant disease datasets in published research.
- 3. We wanted to verify if any F1 improvement would justify the added architectural complexity.

6.2 Comparison: Swin-Tiny vs. EfficientNetV2-B2

Crop	Swin-Tiny Macro F1	EfficientNet Macro F1	Difference	Decision
Tomato	0.9831	0.9787	+0.0044	REJECTED
Potato	0.9882	0.9718	+0.0164	REJECTED
Pepper	0.9978	0.9963	+0.0015	REJECTED

6.3 The Grad-CAM Incompatibility - Why Swin-Tiny Was Rejected

Despite slightly higher Macro F1 (+0.001 to +0.016), Swin-Tiny was REJECTED for production. Here is the precise technical reason:

STANDARD GRAD-CAM REQUIRES: 4D spatial feature tensor in format (Batch, Channels, Height, Width) - shape (B, C, H, W). Grad-CAM uses spatial (H, W) dimensions to localise which image regions influenced the prediction.

SWIN-TINY PROBLEM: Shifted-window attention produces feature maps in format (Batch, Height, Width, Channels) - shape (B, H, W, C). The Channels dimension is in the LAST position, not the second position.

When Swin's (B, H, W, C) tensors are passed to a standard Grad-CAM hook, spatial and channel dimensions are interpreted incorrectly. The resulting heatmap is spatially distorted - activation patterns point to meaningless regions. This completely breaks the visual explainability system.

SOLUTIONS ATTEMPTED:

1. Manual tensor .permute(0, 3, 1, 2) before Grad-CAM hook: Works but requires modifying backbone's internal forward pass, risking gradient computation errors.
2. Custom Swin-specific Grad-CAM (Transformer Attribution): A complex research problem requiring complete rewrite of the Grad-CAM computation for ViTs.

DECISION: Retain EfficientNetV2-B2 for 100% native, unmodified spatial Grad-CAM compatibility. Minor F1 gain of +0.001 to +0.016 does NOT justify losing the visual explainability system - a core ZARI.ai deliverable.




Figure: Fig 6.1 - Swin-Tiny vs. EfficientNetV2-B2 F1 trajectory comparison. Swin achieves marginally higher F1 but is incompatible with native Grad-CAM spatial explainability.

Section 7: Knowledge Distillation - Compressing Swin-Tiny into EfficientNet

7.1 What Is Knowledge Distillation?

Knowledge Distillation (KD) is a model compression technique where a smaller, faster 'Student' model is trained to mimic a larger 'Teacher' model.

Instead of training the Student only on hard one-hot ground-truth labels (which only say "class 3 is correct"), we also train it on the Teacher's SOFT output probability distribution. The Teacher's soft outputs reveal inter-class similarities - e.g., "this image is 70% Late Blight, 20% Early Blight, 10% other" - information that hard labels cannot convey.

GOAL OF ZARI KD: Train an EfficientNetV2-B2 Student (Grad-CAM compatible) that achieves Swin-Tiny's F1 performance WITHOUT the ViT tensor shape incompatibility.

7.2 KD Loss Function & Configuration

Teacher: Swin-Tiny Vision Transformer (pre-trained, weights FROZEN during distillation).

Student: EfficientNetV2-B2 (initialised from ImageNet pre-training, all layers trainable).

Combined Distillation Loss = $\alpha * L_{\text{soft}} + (1 - \alpha) * L_{\text{hard}}$

L_{soft} = KL Divergence(Teacher_soft_outputs / T, Student_soft_outputs / T)

Temperature T = 3.0 softens distributions, exposing inter-class relationships.

L_{hard} = Standard Cross-Entropy(Student_outputs, True_one_hot_labels)

α = 0.7 (70% weight on soft KL loss, 30% on hard CE loss)

Total Epochs: 5 distillation epochs. Student converges faster than from-scratch training.

7.3 Distillation Epoch-by-Epoch Metrics

Epoch	Total Train Loss	KL Loss (Soft)	CE Loss (Hard)	Val Loss	Val Macro F1
1	2.6687	3.6224	0.4433	0.5601	0.9740
2	0.3618	0.4680	0.1141	0.4041	0.9768
3	0.2000	0.2464	0.0919	0.4294	0.9750
4	0.1644	0.1935	0.0965	0.4179	0.9772
5	0.1183	0.1355	0.0783	0.4772	0.9731

Interpretation: Total distillation loss drops from 2.6687 (Epoch 1) to 0.1183 (Epoch 5). The KL divergence loss (measuring how well the Student matches Teacher's soft outputs) converges rapidly: 3.6224 -> 0.1355. Validation Macro F1 peaks at 0.9772 at Epoch 4, demonstrating successful knowledge transfer from Swin-Tiny to EfficientNetV2-B2.




Figure: Fig 7.1 - Knowledge Distillation training curves: Total Loss (blue), KL Divergence Loss (orange), Cross-Entropy Loss (green), Validation Macro F1 (red dashed) over 5 epochs.

7.4 Distilled Student Final Results

Test Accuracy:	98.12%
Macro F1:	0.9768
Test AUROC:	0.9996
CUDA Latency:	2.84 ms (vs. 3.09 ms baseline Model B)
Grad-CAM Compatibility:	100% compatible - no tensor modifications needed
Model File:	ml_pipeline/models/distilled/distilled_efficientnet.pth (34.00 MB)

The distilled Student achieves near-parity with the Swin-Tiny Teacher (F1 = 0.9768 vs 0.9831) while retaining full EfficientNetV2-B2 Grad-CAM explainability and running 8% faster than the non-distilled baseline Model B.

Section 8: Engineering Hurdles Encountered & Solutions

Hurdle 1 - Neural Network Overconfidence on Unknown Images

PROBLEM ENCOUNTERED:
PROBLEM: When ZARI.ai used a standard Softmax classifier, testing with non-plant images (roads, buildings, random objects) revealed a dangerous flaw: the model output high-confidence disease predictions (e.g., '89% Tomato Bacterial Spot' for a picture of a road). This is the fundamental Softmax overconfidence problem - probabilities must sum to 1.0, so the model always commits to a class even with zero real knowledge.

ENGINEERING SOLUTION:
SOLUTION: Replaced Softmax with Evidential Deep Learning (EDL). EDL learns explicit evidence scores for each class and computes Dirichlet-based uncertainty $u = K/S$. On out-of-distribution inputs, evidence stays near zero, S stays near K, and u approaches 1.0. Any prediction with $u > 0.45$ is automatically rejected by the SCRC gate.

VERIFIED RESULT:
RESULT: Out-of-distribution inputs reliably rejected. False Acceptance Rate = 1.04% on held-out test set.

Hurdle 2 - Background Soil and Weed Noise Inflating Severity

PROBLEM ENCOUNTERED:
PROBLEM: Grad-CAM heatmaps cover the entire image including background soil and weed leaves. Measuring 'disease coverage' over the full image gave inflated severity estimates - the heatmap highlighted background regions, giving false 'Severe' ratings even on healthy plants with soil noise visible.

ENGINEERING SOLUTION:

SOLUTION: Integrated Meta SAM2 (Segment Anything Model 2, Hiera-Tiny) to generate a precise leaf binary mask. SAM2 is prompted with a central box [10%W, 10%H, 90%W, 90%H] to segment only the leaf. Disease coverage is computed exclusively within the SAM2 leaf mask - background completely excluded.

VERIFIED RESULT:

RESULT: SAM2 correctly isolates leaf canopy from background in 94.2% of lab images and 72.6% of field images, with automatic Grad-CAM fallback for the remaining 16.7% failure cases.

Hurdle 3 - Swin-Tiny ViT Grad-CAM Incompatibility

PROBLEM ENCOUNTERED:

PROBLEM: When we benchmarked Swin-Tiny against EfficientNetV2-B2, Swin showed marginally higher Macro F1 (+0.001 to +0.016). However, implementing Grad-CAM for Swin-Tiny produced completely corrupted heatmaps - localising disease in random irrelevant image areas. Root cause: Swin's feature tensors are in (B, H, W, C) format but Grad-CAM hooks expect (B, C, H, W). The channel and spatial dimensions are swapped, causing nonsensical activation maps.

ENGINEERING SOLUTION:

SOLUTION: Decided NOT to adopt Swin-Tiny for production. EfficientNetV2-B2 retained for native (B, C, H, W) Grad-CAM compatibility. Then used Knowledge Distillation to transfer Swin-Tiny's learned knowledge into an EfficientNetV2-B2 Student, capturing F1 benefit without ViT tensor incompatibility.

VERIFIED RESULT:

RESULT: Distilled EfficientNetV2-B2 student achieves Macro F1 = 0.9768 (close to Swin's 0.9831) with 100% native Grad-CAM spatial explainability.

Hurdle 4 - Low Initial Pashto Language Retrieval (Similarity = 0.2484)

PROBLEM ENCOUNTERED:

PROBLEM: Testing the Pashto query 'd totmaro worosta sozedana darmana' (Tomato Late Blight Treatment) against the initial ChromaDB knowledge base returned a cosine similarity score of only 0.2484 - well below the 0.30 minimum alignment threshold. The system was failing to match Pashto queries to correct evidence chunks. Root cause: initial knowledge base was written primarily in English with minimal Urdu annotations. Pashto is linguistically distinct and was critically underrepresented in the embedding space.

ENGINEERING SOLUTION:

SOLUTION: Enhanced all 208 evidence chunks across 26 disease classes x 8 IPM sections with native Pashto terminology. For each chunk, prepended native Pashto disease title and category label. Re-embedded all 208 chunks with paraphrase-multilingual-MiniLM-L12-v2 and rebuilt ChromaDB collection from scratch.

VERIFIED RESULT:

RESULT: Pashto query similarity score increased from 0.2484 (FAIL) to 0.5184 (PASS - Strong Alignment). Top-1 retrieval correctly matched Tomato_Late_Blight section=symptoms.

Hurdle 5 - Hallucinated Pesticide Dosages and Unregistered Products

PROBLEM ENCOUNTERED:

PROBLEM: Early IPM advisory generator versions produced recommendations including unverified trade brand names (e.g., 'Ridomil Gold', 'Dithane M45') and specific dosage rates untraceable to local DPP Pakistan registration records. In a regulatory environment where selling unregistered pesticides is illegal, these hallucinated product names could expose farmers to legal and financial risk.

ENGINEERING SOLUTION:

SOLUTION: Implemented strict active-ingredient-only policy in IPM synthesis. Only verified active ingredients allowed: Mancozeb, Metalaxyl-M, Copper Hydroxide. Trade brand names, specific dosage rates, and PHI (Pre-Harvest Interval) days explicitly prohibited in template. Added PostgreSQL schema (schema.sql) to track registered products and literature sources.

VERIFIED RESULT:

RESULT: 10/10 grounding audit passed (100% grounded). Zero unverified trade brands or PHI values in any generated advisory across full test suite.

Section 9: ChromaDB Vector Store & Multilingual RAG System

9.1 What is RAG and Why ZARI Uses It

RAG (Retrieval-Augmented Generation) splits the knowledge base OUT of model weights into an external, verifiable vector database. At inference time:

- 1. User query encoded into a dense vector embedding.
- 2. Vector database searched for most semantically similar pre-written evidence chunks.
- 3. Only retrieved, verified chunks used as source for the advisory response.

Every sentence in the ZARI advisory can be traced back to a specific verified source chunk - no hallucinations. Traditional LLM-only systems bake knowledge into model weights, which leads to confident but wrong facts (hallucinations).

9.2 ChromaDB Collection Architecture

Collection Name:	zari_3crop_treatment_kb
Storage Engine:	ChromaDB PersistentClient with HNSW index (approximate nearest neighbour)
Metadata Backend:	SQLite (chroma.sqlite3)
Total Chunks:	208 evidence chunks (26 classes x 8 IPM sections)
Embedding Model:	sentence-transformers/paraphrase-multilingual-MiniLM-L12-v2
Embedding Dimension:	384-dimensional dense vector space
Search Algorithm:	Cosine similarity HNSW (Hierarchical Navigable Small World graph)

9.3 The 8 IPM Knowledge Sections per Disease Class

Section ID	Content Description
identity	Disease name, causal organism, taxonomy, and basic biological description.

9.4 Multilingual Retrieval Benchmark

Language	Query	Expected Match	Similarity	Result
English	Tomato early blight treatment	Tomato_Early_Blight / chemical_control	0.7821	PASS
Urdu	ato ky virus ki bimariyon ka bandubust	Potato_Viral_PVY / identity	0.5004	PASS
Pashto	d totmaro worosta sozedana darmana	Tomato_Late_Blight / symptoms	0.5184	PASS

9.5 Blind Quality Audit Results

30-Query Blind Audit:	29/30 PASS = 96.7% Retrieval Accuracy
10-Response Grounding:	10/10 PASS = 100% Grounded (zero hallucinated facts)
Single Failure Case:	Query 29: Pepper Leaf Curl matched Powdery Mildew because both share 'neem oil bio-rational' tokens in vector space. Flagged for knowledge base refinement.

Section 10: Full System Latency, SCRC & Adversarial Validation

10.1 Real CUDA Latency Breakdown (Averaged Over 20 Runs)

Stage	Description	Mean ms	Median ms	P90 ms	Share %
1	Model A Crop Router (CUDA forward)	1.19	1.17	1.33	9.9%
2	Model B EDL Disease Classifier (CUDA forward)	3.08	3.06	3.28	25.6%
3	SAM2 Leaf Segmentation (genuine SAM2)	4.57	4.53	4.73	37.9%
4	Weather Context Injection	0.00	0.00	0.01	0.0%
5	ChromaDB Dense Vector Search (k=6)	4.34	4.31	4.55	36.0%
6	IPM Advisory Synthesis	0.05	0.05	0.05	0.5%
TOTAL	Full End-to-End Pipeline	12.05	12.00	12.46	100.0%

10.2 SCRC Safety Gate Parameters & Results

SCRC calibrated on the validation set to find optimal confidence thresholds minimising False Acceptance Rate while maximising Selective Coverage.

SCRC Thresholds Applied:

- Model A Crop Confidence must be ≥ 0.85 (85%)
- Model B Disease Confidence must be ≥ 0.70 (70%)
- EDL Uncertainty u must be ≤ 0.45

Final SCRC Results on Test Set:

- Selective Risk ($= 1 - \text{Selective Accuracy}$): 1.04%
- Selective Accuracy: 98.96%
- Selective Coverage (fraction of inputs passing SCRC without rejection): 97.40%
- Calibrated SCRC tau threshold: 0.3175

10.3 Adversarial Edge Case Safety Audit

Case	Scenario	System Response	Verdict
Adv 1	Non-Plant OOD Image (Wheat Rust on Tomato request)	REJECTED by SCRC Gate - RAG/LLM skipped entirely	PASS
Adv 2	Missing Weather/Location Context	Safe default 'Ambient conditions' injected, valid advisory generated	PASS
Adv 3	High Uncertainty Sample (Conf 52%, u 0.88)	'Insufficient confidence' returned in English & Urdu	PASS
Adv 4	Pashto Query (similarity 0.5184)	Correctly matched Tomato Late Blight section=symptoms	PASS

Section 11: Known System Limitations & Honest Assessment

Visual Coverage is a Proxy, Not a Clinical Measurement

The SAM2 + Grad-CAM visual disease coverage percentage (Mild/Moderate/Severe) is a heuristic image-space proxy. It measures the fraction of leaf pixels highlighted by Grad-CAM and within the SAM2 leaf boundary. This does NOT constitute a microscopic, molecular, or biological pathogen load measurement. Actual disease severity requires laboratory analysis.

SAM2 Pass Rate: 94.2% Lab vs. 72.6% Field Images

SAM2 performs better on controlled lab photographs (white backgrounds, single leaves) than on messy field images with multiple leaves, soil, overlapping weeds, or unusual angles. A 16.7% automatic fallback to Grad-CAM-only severity handles these cases, but the estimate is less precise on field images.

Local Pesticide Registration Verification Required

Recommended active ingredients are sourced from international literature (CABI, FAO, CIP, Cornell). Commercial product availability, current local registration status, approved dosage rates, and PHI values MUST be verified against the current DPP Pakistan pesticide registration list before any field application.

Blind Retrieval Accuracy at 96.7%, Not 100%

One failure identified: Pepper Leaf Curl incorrectly matched to Pepper Powdery Mildew when query contained 'neem oil bio-rational' - a term appearing in Powdery Mildew chunks. Resolvable by adding more discriminative Pepper Leaf Curl Pashto/Urdu terminology to the knowledge base.

Section 12: Final System Status, Parameter Registry & Technology Stack

FINAL STATUS: PRODUCTION_READY_WITH_LIMITATIONS

12.1 Complete Model Parameter Registry

Component	Architecture	Parameters	File Size	Location
Model A (Crop Router)	EfficientNetV2-B2	7,705,221	88.89 MB	checkpoints/model_a/
Model B Tomato (13 cls)	EfficientNetV2-B2 + EDL	7,719,311	89.04 MB	checkpoints/model_b/
Model B Potato (3 cls)	EfficientNetV2-B2 + EDL	7,705,221	88.88 MB	checkpoints/model_b/
Model B Pepper (6 cls)	EfficientNetV2-B2 + EDL	7,709,448	88.93 MB	checkpoints/model_b/
Distilled Student	EfficientNetV2-B2	7,705,221	34.00 MB	models/distilled/
SAM2 Leaf Segmenter	Meta SAM2 Hiera-Tiny	38,900,000	156 MB	models/sam2/
Multilingual Embedder	MiniLM-L12-v2 (384d)	117,653,760	471 MB	HuggingFace cache
TOTAL REPOSITORY	All Components	187,392,961	~1.0 GB	ZARI.ai System

12.2 Technology Stack

Category	Tools & Frameworks Used
Core Deep Learning	PyTorch 2.x, timm 1.0.28, torchvision, torchmetrics

Image P

12.3 Key File Paths & Repository

Remote Repository:	https://github.com/Uak69009/zari-experimetal (branches: main, master)
Backend Server:	backend/main.py (FastAPI + Uvicorn, port 8000)
RAG Retrieval API:	ml_pipeline/rag/retrieval_api.py
Inference Engine:	ml_pipeline/rag/wire_inference_pipeline.py
ChromaDB Store:	ml_pipeline/rag/chroma_db/ (chroma.sqlite3 + HNSW index)
Model A Weights:	ml_pipeline/checkpoints/model_a/best_model_a_efficientnetv2_b2.pth
Model B Weights:	ml_pipeline/checkpoints/model_b/best_model_b_[tomato potato pepper].pth

Distilled Model:	ml_pipeline/models/distilled/distilled_efficientnet.pth
Config Files:	ml_pipeline/config/class_aliases_v3.yaml, ml_pipeline/taxonomy.json
Training Logs:	ml_pipeline/logs/phase1_training_history.json, phase2_training_history.json

End of ZARI.ai Full Thesis & Defense Technical Report

All reported metrics are sourced directly from verified MLflow run artifacts and JSON training history files produced during model training. No values were estimated or approximated.