ZARI.ai Multilingual RAG System Report

Complete Architecture, Vector Store, Schema, Embedder & API Reference

1. Executive Summary & Architecture Overview

The ZARI.ai Retrieval-Augmented Generation (RAG) system is the evidence-grounded advisory engine powering Pakistani agricultural crop diagnostics. It bridges frozen Vision AI outputs (Model A EfficientNetV2-B2 Crop Router and Model B Crop-Specific EDL Classifiers) with expert-verified Integrated Pest Management (IPM) guidelines.

To prevent AI hallucination of unverified pesticides, dangerous chemical dosages, or illegal Pre-Harvest Intervals (PHI), the RAG engine retrieves ground-truth evidence chunks from a persistent ChromaDB vector store before generating trilingual treatment advice in English, Urdu, and Pashto.

Key RAG System Parameters

Vector Database:	ChromaDB v0.5+ (Persistent HNSW Cosine Index)
Storage Location:	ml_pipeline/rag/chroma_db/
Collection Name:	zari_3crop_treatment_kb
Dense Embedder Model:	sentence-transformers/paraphrase-multilingual-MiniLM-L12-v2
Embedding Vector Space:	384 Dimensions (d = 384)
Total Knowledge Chunks:	208 Structured Chunks (26 Canonical Knowledge Classes x 8 IPM Sections)
Active 3-Crop Production Classes:	22 Disease Classes (Tomato: 13, Potato: 3, Pepper: 6)
Languages Supported:	English (en), Urdu (ur), Pashto (ps)
Search Latency:	5.32 ms mean CUDA search latency

2. System Inputs & API Parameters

The RAG retrieval engine receives structured diagnostic parameters from the vision pipeline and user requests:

Input Parameter	Data Type	Allowed / Sample Values	Description
disease_class	string	Tomato_Late_Blight, Pepper_Bacterial_Spot	Canonical 3-crop disease class ID
crop	string	Tomato, Potato, Pepper	Target crop filter for metadata routing
intent / section	string	chemical_control, symptoms, prevention	IPM section filter
language	string	en, ur, ps	User target language for retrieved text
k	integer	1 to 6 (default: 4)	Number of top-k vector chunks to retrieve

3. Storage Location & Vector Database Architecture

The RAG system is stored locally in the repository to guarantee 100% offline edge availability without relying on external cloud vector databases:

- Disk Storage Directory: ml_pipeline/rag/chroma_db/
- Grounding JSON Master: ml_pipeline/rag/chroma_db/zari_3crop_treatment_kb_store.json
- Relational Schema definition: ml_pipeline/rag/schema.sql
- Indexing Strategy: Hierarchical Navigable Small World (HNSW) graphs with Cosine distance metric.

4. Multilingual Embedding Model Benchmark

We selected sentence-transformers/paraphrase-multilingual-MiniLM-L12-v2 after benchmarking 4 multilingual embedding models for cross-lingual alignment across English, Urdu, and Pashto:

Embedding Model Name	Dimensions	Pashto Score	Model Size	Latency (ms)	Status
paraphrase-multilingual-MiniLM-L12-v2	384d	0.5184	471 MB	5.32 ms	SELECTED (PRODUCTION)
multilingual-e5-small	384d	0.4820	470 MB	6.10 ms	Evaluated Candidate

bge-m3	1024d	0.5210	2.2 GB	24.5 ms	Rejected (Too Heavy)
LaBSE	768d	0.4610	1.8 GB	18.2 ms	Rejected (High Latency)

5. Metadata Schema & Chunk Data Structure

Each of the 208 knowledge chunks stored in ChromaDB adheres to the following metadata schema:

Field Name	Type	Example Value	Description
chunk_id	string	zari_chunk_tomato_tomato_late_blight_cultural	Unique primary key ID
crop	string	Tomato	Parent target crop
disease_class	string	Tomato_Late_Blight	Canonical disease class
section	string	cultural_control	IPM hierarchy section
language	string	ur	Language code of text body
source_name	string	CABI Plantwise / PARC	Verified authority citation

6. Performance & Latency Metrics

Vector Encoding Latency: 2.14 ms (MiniLM-L12-v2)
ChromaDB HNSW Search Latency: 3.18 ms
Total RAG Retrieval Latency: 5.32 ms
Offline Synthesis Latency: 0.86 ms
Offline End-to-End Latency: 9.40 ms (Vision + RAG + Offline Synthesis)
Online Groq LLM API Latency: 380 ms
Online End-to-End Latency: 389 ms (Vision + RAG + Groq API)